


Intelligent Information Systems

SS 2010

2. Deductive Databases and Datalog



- | | |
|------------|--------------------------------|
| 2.1 | Deductive DB: Overview |
| 2.2 | Datalog: Learning by Doing |
| 2.3 | Datalog: Facts and Rules |
| 2.4 | DDL and DML for Datalog |



- 2.1 Deductive DB: Overview
- 2.2 Datalog: Learning by Doing
- 2.3 Datalog: Facts and Rules
- 2.4 DDL and DML for Datalog**



2.4.1 Data Definition Language

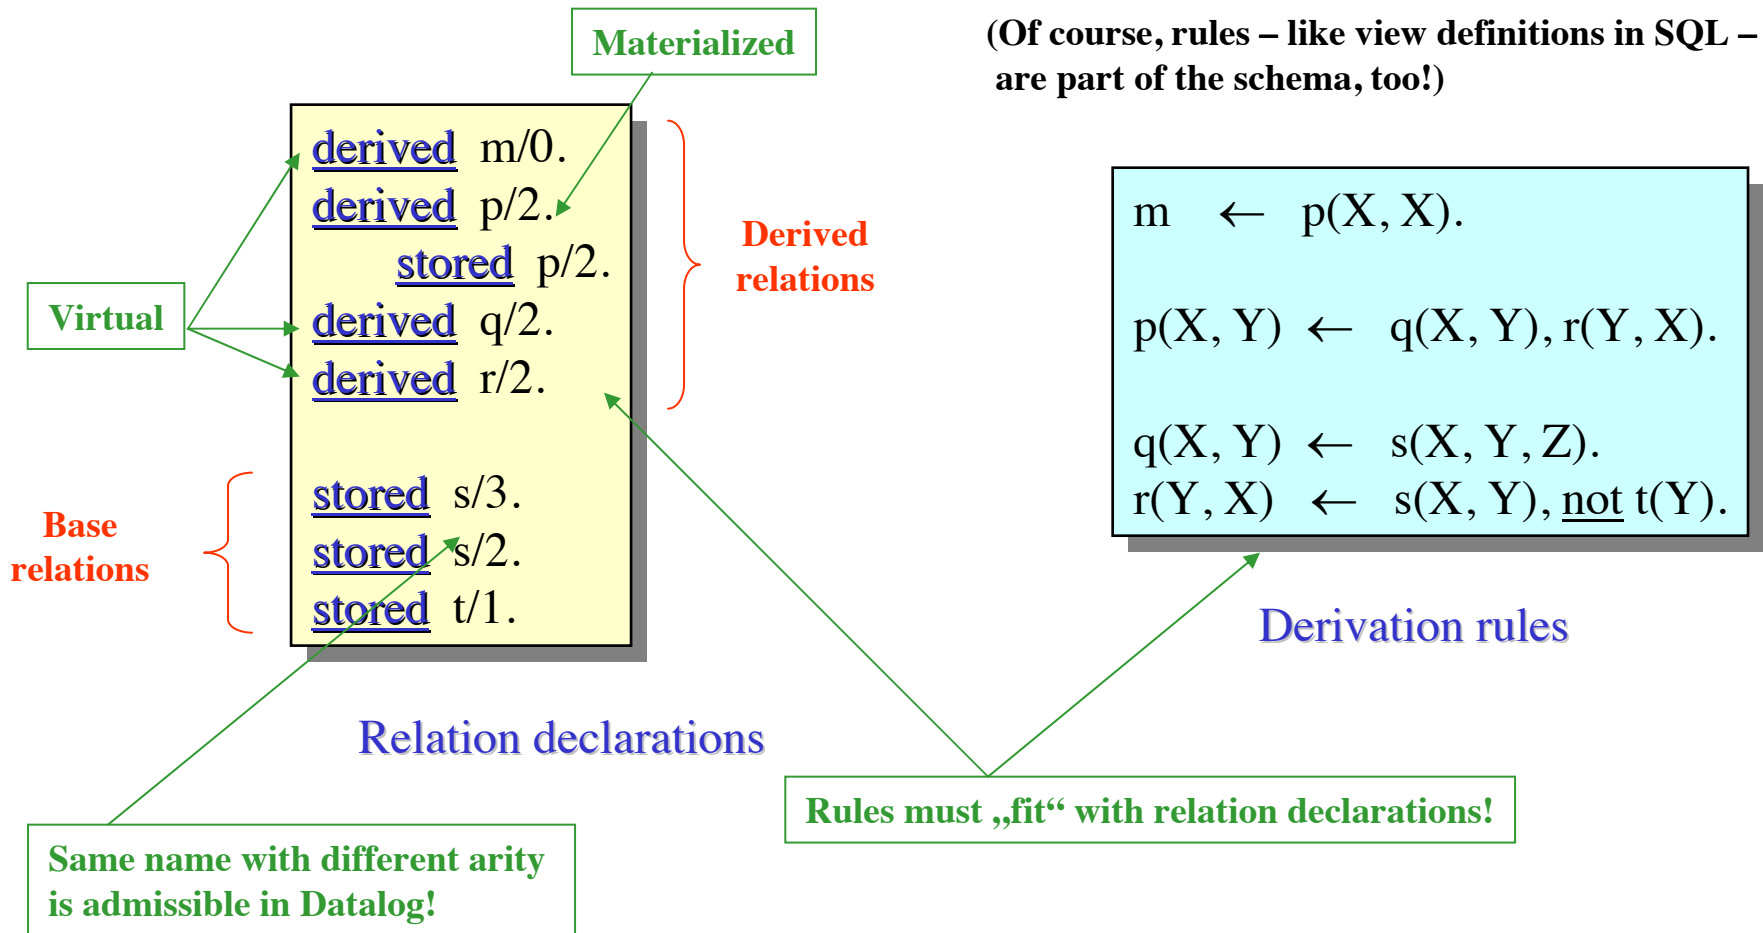
- As for each database: **Declaration** of every relation is required, before the DB can be opened and facts be inserted (or rules be defined).

Database Schema

- In Datalog: There is **no agreed notation** or convention for schemas; most authors simply don't mention a schema explicitly – we have our own notation.
- Datalog is an **untyped language** (like Prolog), i.e., there are no attributes and no explicit value domains. However, we will use types implicitly throughout.
- In a schema: Only the **arity** of a relation and its **name** are declared.
- The **"form of definition"** of the relation, however, is to be fixed in the schema, too:
 - Relation defined by facts only: stored
 - Relation defined by rules only: derived
- Both variants can be combined, i.e., there are three **relation types**:
 - **Base relation**: stored (but not derived)
 - **virtual relation**: derived (but not stored)
 - **materialized relation**: both, stored and derived

engl. "to store": **speichern**
engl. "to derive": **ableiten**

Example of a simple **Datalog** schema:



Surprisingly, there is **no agreed notation for integrity constraints** in Datalog either – thus, we introduce our own style for expressing normative rules, too:

derived m/0.
derived p/2.
 stored p/2.
derived q/2.
derived r/2.

stored s/3.
stored s/2.
stored t/1.

m ← p(X, X).
 p(X, Y) ← q(X, Y), r(Y, X).
 q(X, Y) ← s(X, Y, Z).
 r(Y, X) ← s(X, Y), not t(Y).

Constraints may apply to both,
 stored and derived relations.

Positive condition:
 „Must always be true!"

Negative condition:
 „May never be true!"

constraint t(c).
constraint not m.
constraint not q(a, b).

Integrity constraints

In our sample database representing a family tree a lot of „nonsense“ can be accomodated without integrity constraints, e.g.:

born_in('John', 1978).
born_in('Mary', 1933).
born_in('Mary', 1968).

died_in('Jim', 1999).
died_in('John', 1945).
died_in('Jim', 2000).

sex_of('Jim', n).
sex_of('male', m).

child_of('William', 'Diana', 'Charles').
child_of('John', 'Caroline', 'William').
child_of('Charles', 'Christine', 'John').

married_to('Charles', 'Diana', 1981).
married_to('Charles', 'Anne', 2004).

divorced_from('Diana', 'Charles', 1996).
divorced_from('Elizabeth', 'Charles', 2002).

title_of('Charles', 'King of Great Britain').
title_of('William', 'King of Great Britain').

Each of these facts is somehow "inadmissible"! How to prevent them in Datalog?

ICs for the genealogy DB: (formulated in natural language first)

- IC 1 Each person who has died must have been born before.
- IC 2 There is just one person per name.
- IC 3 Each person's sex is (either) male or female.
- IC 4 If two people divorce they have to be married at the time of divorce.
- IC 5 In the 'married_to'-relation, the 1st column always contains the name of the wife, whereas the husband is recorded in the 2nd column.
- IC 6 (analogously for 'divorced_from')
- IC 7 (analogously for 'child_of')
- IC 8 The 1st column in the 'sex_of'-relation contains the name of a person.
- IC 9 Nobody must be a descendant of himself or of one of his own descendants (i.e., the 'descendant_of'-relation is acyclic).
- IC 10 Nobody must be married to a relative of degree 1 or 2.
- IC 11 Each title must be held by at most one living person a time.
- IC 12 Nobody may get married before being born.

IC 1 Each person who has died must have been born before.

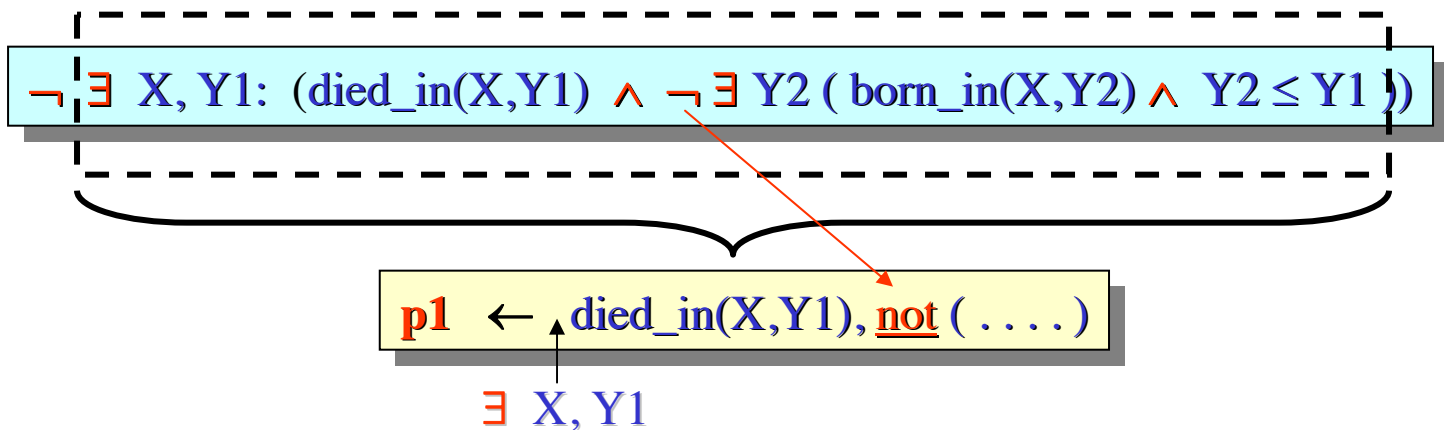
- Most integrity constraints are **universal** statements (like IC1):
 "For all" objects of a particular type a certain condition holds!
- "For all" is circumscribed in natural language most of the time (even negatively):
 "each", "only", "all", "always", "nobody"
 (or is not made explicit at all, as in IC 4 and IC 8)
- In (full) predicate logic, such conditions are expressed using a universal quantifier, in most cases coupled with implication for restricting the scope of the statement:

$$\forall X, Y1: \text{died_in}(X, Y1) \Rightarrow \exists Y2 (\text{born_in}(X, Y2) \wedge Y2 \leq Y1)$$


- Like SQL, Datalog does not offer any universal quantifier, but only the (implicit) existential quantifier, so that a transformation using Boolean algebra is required:

$$\neg \exists X, Y1: (\text{died_in}(X, Y1) \wedge \neg \exists Y2 (\text{born_in}(X, Y2) \wedge Y2 \leq Y1))$$


- Universal conditions are thus turned into **negated existential conditions**.
- Existential conditions in Datalog are expressible via **parameterless relations** (arity 0) defined (once again – a very important concept!) by auxiliary rules:



- As nesting is banned from Datalog, the conjunction has to be moved "behind" the negation operator and pushed into another auxiliary rule:



Thus, the **positive part** of the existential condition has been properly formulated in Datalog:



What is missing is just the integrity constraint itself (including the negation), which we will identify from now on by the keyword constraint:

constraint $\text{not } p1$

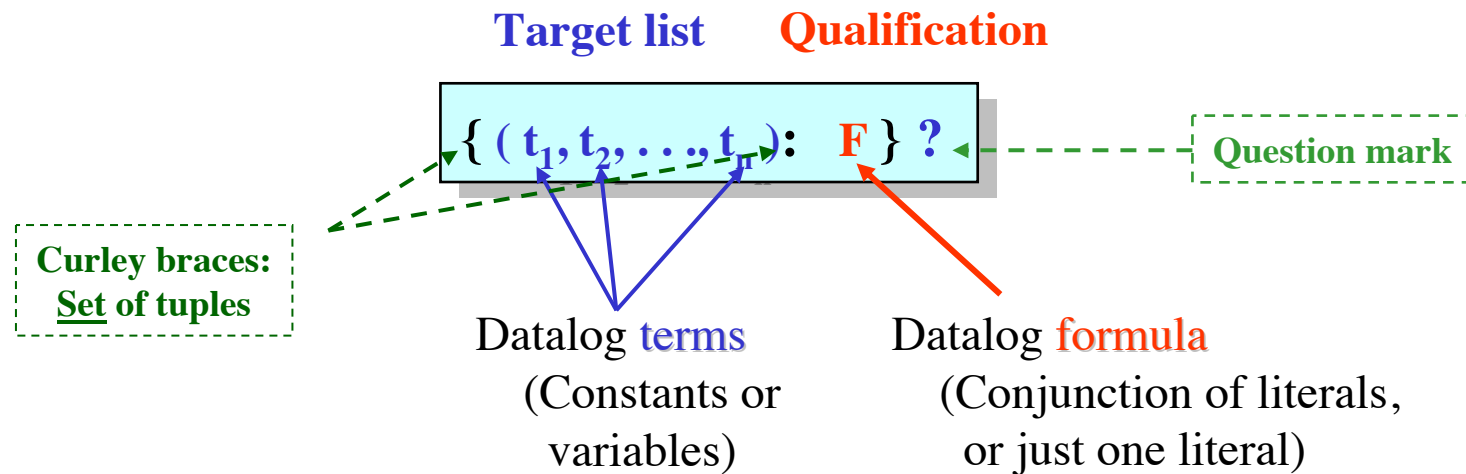
- A **Datalog schema** consists of the following parts:
 - **Relation declarations** of the form $\text{[derived | stored] } \langle \text{Name} \rangle / \langle \text{Arity} \rangle$
 - Relation names are strings beginning with a lower case character.
 - Arities are positive integers or zero.
 - Per couple name/arity occurring in the DB there must be at least one declaration; double declarations (derived and stored) are possible.
 - **Derivation rules** of the form $\langle \text{Head} \rangle \leftarrow \langle \text{Body} \rangle.$
 - The rule head is a positive literal for a relation declared as derived.
 - The rule head is a conjunction of literals (or just a single literal).
 - Each derived relation is defined by at least one rule.
 - **Integrity constraints** of the form $\text{constraint } \langle \text{Literal} \rangle.$
 - The literal must be variable-free (a ground literal, positive or negative).
- In addition, various "wellformedness conditions" have to be respected.



- 2.1 Deductive DB: Overview
- 2.2 Datalog: Learning by Doing
- 2.3 Datalog: Facts and Rules
- 2.4 DDL and DML for Datalog**

2.4.2 Data Manipulation Language

- As in each query language, there are in Datalog **two types of queries** in Datalog, too:
 - Retrieval queries** serve as a means of access to facts in the DB.
 - Test queries** are employed if testing yes/no conditions.
- The syntax of **retrieval queries** closely resembles the mathematical concept of an intensional set expression $\{X \mid f(X)\}$ ("Set of all x, such that f(X) holds")



Examples of **retrieval queries** (with their corresponding answer sets):

$\{ (S) : \text{son_of}(S, \text{'Elizabeth'}) \} ?$

Which sons does Elizabeth have?

Answer: $\{ (\text{'Charles'}), (\text{'Andrew'}), (\text{'Edward'}) \}$

$\{ (X) : \text{grandchild_of}(X, \text{'Elizabeth'}), \text{mother_of}(Y, X), \text{not divorced}(Y) \} ?$

Which grandchildren of Elizabeth have a mother, who is presently not divorced?

Answer : $\{ (\text{'Peter'}), (\text{'Zara'}) \}$

Apparently: Queries and rules are syntactically **very similar!**

Target list	\Leftrightarrow	Rule head
Qualification	\Leftrightarrow	Rule body

e.g.:

$\{ (X, Y) : \text{father_of}(X, Z), \text{child_of}(Y, Z) \} ?$

$\text{grandfather_of}(X, Y) \leftarrow \text{father_of}(X, Z), \text{child_of}(Y, Z) .$



Derivable relation:

Named answer set of a query

- Due to this analogy

Safety requirement for queries necessary!

- "Safe queries":
 - All variables in the **target list** must occur in the qualification, too.
 - All variables in **negative literals** must occur in at least one positive literal, too.
- In analogy with rules, too:
 - **Local variables** in the qualification are implicitly **existentially quantified** at the beginning of the qualification part.
 - Answer sets are **free of duplicates**.

Rules are "stored queries".

- Various reasonable queries **cannot be expressed** up till now, due to syntactic restrictions in Datalog such as forbidden nesting and missing disjunctions, e.g.:

Find all persons who are **first or second cousins** of William?

- Queries like this require **rules** to be added, even though these rules are not really suitable as parts of the DB schema (being rather occasional or „ad hoc“ in nature) – in our example:

```
cousin1_or_2_of (X,Y) ← cousin_of(X,Y).
cousin1_or_2_of (X,Y) ← cousin_of(X,Z), parent_of(Z,Y).
```

**Implicit
disjunction**

- The query itself may then refer to the new rule-defined concept:

{(X) : **cousin1_or_2_of** (X, 'William')} ?

- As such rules better don't „pollute“ the schema, it is a good idea to „embed“ them into the respective query, indicating that they will not be retained outside this query:.

⇒ **"Local Rules"** inside a query are in fact inevitable !

- In the example:

```
{ (X) : cousin1_or_2_of (X, "William") }
      ( with , cousin1_or_2_of (X,Y) ← cousin_of(X,Y) ;
        cousin1_or_2_of (X,Y) ← cousin_of(X,Z), parent_of(Z,Y) ?
```

Additional delimiter for
lists of local rules!

- The newly introduced keyword with for indicating that rules are local to a query has been inspired by the same word in SQL:1999, where it serves a similar purpose within view declarations.
- Locally defined relations may not be used outside the respective query and are no part of the DB schema!

- Syntactically admissible, but rarely necessary: **Constants in the target list**

e.g.:

$$\{ (X, \mathbf{a}, Y) : p(X, Y) \} ?$$

$$\{ (\mathbf{a}) : p(X, a) \} ?$$

- Finally as a kind of "extreme case": Queries with an **empty target list**,
a funny way of expressing Boolean existential queries

$$\exists X :$$

$$\{ () : p(X, a) \} ?$$

Antwort	$\{ () \}$	\longrightarrow	<u>true</u>
	$\{ \}$	\longrightarrow	<u>false</u>

But that's the way
you have to do it
in SQL, folks!

- Rather than forcing users to formulate test queries in this style as "degenerate" retrieval queries (as in SQL), „our“ Datalog offers a **special syntax** for this:
 - based on **parameterless rule heads** (implicit \exists),
 - or consisting of **ground literals** only (or both)

exists_widower
with exists_widower \leftarrow widower(X) ?
single('Charles') ?

$\exists X :$

- **Negated** or **conjunctive** test queries are possible, too :

single('Charles'), widower('Charles') ?
not exists_orphan
with exists_orphan \leftarrow orphan(X) ?

- Important: It is not sufficient to introduce a parameterless relation defined by a derivation rule for expressing a constraint!
- Reason: Parameterless relations may be queried (as tests), but are not automatically combined with the request to the DBMS to automatically enforce this test!

1) derived consistent/0.

consistent \leftarrow p(X), not q(X).

2) derived consistent/0.

constraint consistent.

consistent \leftarrow p(X), not q(X).

3) constraint consistent

with consistent \leftarrow p(X), not q(X) .

1) and 2): Test query "consistent ?"
can be posed

3): Relation "consistent/0" is not queriable
(as local to the integrity constraint)

1): Answer may be true or false.

2): Answer is always true.

- For **updates** (DB changes, i.e., insertions, deletions and modifications of facts) in Datalog there is no generally accepted syntax and semantics in the literature either. Again we need a notation of our own, which seamlessly fits with the rest of the language. (Note that we don't consider schema changes in this lecture.)
- We will introduce **insertions** (+) und **deletions** (−) only:
Modifications of individual attribute values can be simulated by combining +/−.
- For now: **Only updates of base relations are considered !!**
- Syntactical buliding block of all Datalog updates:
 (Literals with a sign + / −)
- All update statements are delimited by an (imperative) **exclamation mark**.

Dynamic Literal

e.g.:

```
+ successor('Elizabeth', 'William') !
− successor('Elizabeth', 'Charles') !
```

- **Elementary updates** (insertions/deletions) of individual facts in a relation are expressed by means of dynamic ground literals, e.g.:

+ title('Paul McCartney', 'Sir') !

- **Conditional updates**: Several, simultaneously executable elementary updates depending on a condition (syntactically like a rule body)

+ title(X, 'Sir') : member_of(X, 'Beatles') !


Dynamic literal


Condition

- **Safety requirement** for conditional updates: In analogy to rules and queries
- For conditional updates, local rules are possible, too:

+ p(X) : r(X), not s(X)
with s(X) ← t (X,Y) !

- The analogy already observed for queries and rules carries over to conditional updates:

dynamic literal	\Leftrightarrow	target list	\Leftrightarrow	rule head
condition	\Leftrightarrow	qualification	\Leftrightarrow	rule body

e.g.:

{ (X, Y) : father_of(X, Z), child_of(Y, Z) } ?

grandfather_of(X, Y) \leftarrow father_of(X, Z), child_of(Y, Z) .

+ grandfather_of(X, Y) : father_of(X, Z), child_of(Y, Z) !


- For elementary updates/facts/test queries similarly:

+ born_in('Paul', 1943) !

born_in('Paul', 1943) .


born_in('Paul', 1943) ?

- Conditional updates are **interpreted in a set-oriented** manner (so are retrieval queries and derivation rules!):




- All p-insertions are performed simultaneously.
- The condition is completely evaluated before any change takes place.


- Condition evaluation and DB-update are performed strictly separately **in two phases**. Thus, an update may not cause any **retroactive effects** on the evaluation of its own condition:



- Elementary updates which don't have any changing effect on the current state are accepted: **Redundant updates**



- This is useful in particular for conditional updates where some of the variable bindings resulting from condition evaluation have an effect, whereas others don't:



- **Transactions** for databases are " **blocks of changes** " either executed entirely, or not at all.

⇒ State transitions between DB states are expressible via transactions which could not be executed individually due to integrity violations.

- In SQL: Transactions are sequences, introduced and concluded by special keywords.
- For Datalog: "**Set-oriented**" **transaction concept** seems to be appropriate, i.e.:

Simultaneous execution of several changes without fixing any ordering of steps.

- **Syntax format** for transactions:

$\{ U_1; U_2; \dots; U_n \} !$

(U_i elementary or conditional updates without '!')

- If using conditional updates within a transaction, it may happen that **contradictory** elementary updates result from evaluating different parts of the transaction. Here, *contradictory* means that the same elementary update appears with positive and negative sign, e.g.:

$\{-p(a); +p(X) : q(X)\} !$

q(a). p(a).
q(b). p(c).
q(d).


- Datalog transactions do not fix any ordering of execution, but are set-oriented (in order to remain compatible with other Datalog conventions). Therefore, it is necessary to eliminate elementary updates mutually compensating each other. We call this step "**net-effect**" computation

~~-p(a)~~
~~+p(a)~~
+p(b)
+p(d)



q(a). p(a).
q(b). **p(b).**
q(d). p(c).
 p(d).

- It may happen that different parts of a transaction appear to cause the **same elementary update** to be executed several times:



- This kind of "redundancy" among dynamic literals is to be treated according to the general design principles of Datalog, too. As duplicates are eliminated in queries and during rule application (due to sets being free of duplicates), we will eliminate all dynamic literals appearing more than once before computing contradictory updates and net effect.
- Common, recurring principle underlying all these conventions:
 - No duplicates** in Datalog (all results are sets).
 - No ordering** in Datalog.

- Datalog's DML offers syntactical means for expressing **queries** and **updates**.
- **Queries** are either ...
 - ... **retrieval queries** in the format { <Target list> | <Qualification> } ...
 - ... or variable-free **test queries** in the format <Qualification>
(Qualifications are literals or conjunctions of literals.)
 - All queries are terminated by the symbol **?**.
- **Updates** are either ...
 - ... **dynamic literals** in the format + <Literal> or - <Literal> ...
 - ... or **transactions** in the format {<List of updates>}.
 - All updates are terminated by the symbol **!**.
- Queries and updates have to be **safe** – like rules.
- Queries and updates have a **set-oriented semantics**.
- Updates are executed only if integrity constraints are not violated after the change.